

02

The Attention Devices: Expressive Subjects and the Moment of Use

Document 2 of 8

CONTENTS

1 The Face Gets the First Glance — Then What?

2 The Gaze Rule: A Face Must Have a Job

3 Expression Is a Claim You're Making

4 Craft Rules at the Size That Matters

5 The Moment of Use: Action Without a Face

6 One Subject, Always — and the Side-by-Side Exception

7 Dialing the Devices Across the Funnel

8 Key Takeaways

The Face Gets the First Glance — Then What?

There is a part of your reader's brain that finds faces before the rest of the brain has decided to pay attention to anything. It is not a preference; it is wiring. Eye-tracking research keeps finding the same thing across print, banners, and feeds: a clear human face pulls the very first fixation, ahead of conscious thought. And the pull belongs to expressive faces generally, not to human ones alone — a dog's cocked-head, ears-up expression is read just as fast and just as emotionally as a human face. In a native feed stuffed with editorial-looking photos, an expressive face is the most reliable scroll-stopper you have — and it costs you nothing, because a candid shot of a face mid-feeling is exactly the kind of image the feed already expects to contain.

Lead with an expressive subject caught in a legible moment. The attention device is whatever carries the story's emotion or action at a glance — a face with a natural, readable expression, or the moment of use (mid-action). Crop tight enough that the emotion or action is unmistakable at phone size.

So faces work. Here is the trap: capturing the glance is only half the job, and it's the easy half. A face that grabs the eye and then does nothing with it is attention-stealing decoration — it stops the scroll, delivers the reader nowhere, and gets scrolled past a half-second later. Most weak ad images aren't weak because they lack a face. They're weak because someone dropped a face in *knowing faces work* without deciding what the face was supposed to accomplish. This document is about making the attention device in your image earn its place: the expressive subject, what its gaze is doing, what its expression is promising — and the equal peer of the face, the moment of use: the subject mid-action.

The Gaze Rule: A Face Must Have a Job

What decides whether a face pays off is almost entirely about the eyes. Where a face is looking tells the reader what to do next, and there are exactly three things a gaze can do.

Direct gaze — the eyes meet the reader. When the subject in your image looks straight down the lens, the reader feels addressed. It's the visual equivalent of someone turning to face you and starting to talk. This creates connection: *this is a person like me, and this is*

about people like me. Reach for direct gaze when the image's job is relatability and trust — testimonial-flavored moments, the relatable-problem callout, the "someone in your situation" framing. The eyes say *you*.

Aimed gaze — the eyes look at something in the frame. This is the workhorse, and it's the most underused. When a face looks at an object inside the image, the reader's eye follows the gaze straight to that object. Eye-tracking work has shown this again and again: a face gazing toward a thing pulls fixations onto that thing and measurably lifts the reader's interest in it. The face becomes an arrow made of a living being — a cat staring fixedly at a half-hidden toy transfers the reader's attention to that toy exactly the way a person's gaze would. Use aimed gaze whenever there's a subject you want the reader to end on — the person looking at the thing the story is about, so the reader ends up looking at it too.

Jobless gaze — the eyes point at nothing. This is the mistake. A face staring off-frame at empty space, or a face pasted into a composition with no relationship to anything else in it, captures the reader's attention and then squanders it. The eye lands, finds no direction, and leaves. A jobless gaze is worse than no face at all, because it spent your one fixation and gave nothing back. Before you use a face, answer the question the reader's eye is about to ask: *where do I look now?* If the face doesn't answer it — through connection or through direction — the face isn't finished.

Expression Is a Claim You're Making

A face doesn't just point. It emotes, and its emotion is a promise about the story behind the click. A worried face promises a problem story. A relieved face promises a resolution story. An amazed face promises a discovery story. The reader reads that promise instantly, before a single word of the headline, and the promise has to match the angle the headline is actually carrying. Put an amazed, jaw-dropped expression over a calm, practical headline and the two contradict each other — the reader senses the seam, reads the whole unit as staged, and discounts it. The image and the headline are processed as one object, so a mismatched expression doesn't just fail on its own; it drags the words down with it.

One warning from the research worth internalizing: not every reader reads an expression the same way, and the expressions most likely to be *misread* are the exaggerated, theatrical ones. The over-acted grin, the cartoon shock — these are the faces that scan as

fake or land as the wrong emotion entirely. Natural, recognizable emotion in a real context beats acted intensity every time. You want the expression a real person actually wears in that moment, not the one a stock model performs for a camera. The rule holds whatever the subject: the expression must be natural and readable.

Craft Rules at the Size That Matters

All of this has to survive being shrunk to a few hundred pixels and glanced at mid-scroll on a phone. That reality dictates the craft.

Crop to the focal point. Crop to whatever carries the emotion or action — the expression, the gesture, the moment — tight enough that it's unmistakable at feed size. A full-body shot turns the expression into an unreadable smudge at that scale. How tight scales with where the image lives: on the ad, crop tightest, so the focal point fills the frame; on the advertorial and jump pages, surrounding context is allowed as long as the focal point still reads first. When in doubt, crop tighter.

One subject, one emotional read. Two or more subjects split that precious first fixation and blur the emotional read — the reader doesn't know whose feeling to catch. Group shots belong to brand advertising, which is a different game with different goals. In this system, one subject, one feeling, one clear read.

Real over posed. A candid, slightly imperfect face in a real setting belongs in the feed. A polished stock-model grin announces "advertisement" and gets skipped. Authentic beats produced — this is the register the whole system runs in, and it's the subject of *Document 5, The UGC Default*.

Never a real, identifiable person — and never a celebrity or public figure, full stop. That second one isn't a gray area; it's an outright ban (see *Document 7, The Seven Bans*). AI-generated faces are fine, with one habit attached: inspect them for artifacts — the odd extra tooth, the melted ear, the wrong number of knuckles — before anything ships.

The Moment of Use: Action Without a Face

A face is not the only way to put a subject in the frame, and sometimes it's the wrong way. The broader device is the **moment of use** — the subject mid-action — and its most common instance is hands doing something: gripping, opening, applying, mid-task. This is the second of the two attention devices — a full peer of the face, not a fallback — and it solves problems a face can't. A dog mid-leap catching the toy is the moment of use — no hands required.

Start with what a hand does to an object. An object sitting alone on a table is a product shot: static, anonymous, ad-like. Drop that same object into a hand and it becomes a *moment* — someone is using this, right now, and the reader's mind reflexively completes the story of what's happening. The hand converts an object into an action. That conversion is why hands earn their place, and it works for three reasons.

Hands demonstrate instead of display. A hand mid-action answers the reader's silent first question — *what is this and what do you do with it?* — in a single glance, before any words are read. A product shot makes the reader work that out. A hand shows them.

Hands read as UGC by default. A hand-held shot implies a phone in the *other* hand — it looks like a photo a real person snapped of their own life, which is precisely the register the feed rewards. Object-on-a-white-table implies a studio and a marketer. Hand-in-a-real-setting implies a person. (More on why that authenticity is non-negotiable in *Document 5, The UGC Default*.)

First-person POV puts the reader inside the scene. Shoot from the owner's own perspective — the reader's hands, doing the thing — and the image stops being about someone else and becomes a preview of the reader's own experience. First-person point of view is the most immersive framing you can get without a face, and it quietly sidesteps a pile of casting decisions: there's no expression to match, no face for the reader to relate to or fail to relate to, no likeness to worry about. When the story is about the doing — the mechanism, the ritual, the moment of use — first-person hands are often the strongest shot available.

The craft rules for hand-led images follow the same feed-size logic as faces. Crop tight enough that the action is unmistakable at phone size. One action, not a flurry of several. Real skin, real setting, natural light — a manicured hand on a seamless white sweep is a

stock photo wearing a disguise, and the feed sees through it. And here's the action-device version of the jobless gaze: the action must be *the story's* action. A hand doing something unrelated to your angle wastes the fixation exactly the way a face staring at nothing does.

So choose the device by what the story is about: an expressive face — human or animal — when the story is the feeling; the moment of use when the story is the doing.

One Subject, Always — and the Side-by-Side Exception

Everything above assumes a discipline that underlies all ad-image composition: **one clear focal subject**. One subject, one moment. The eye gets a single fixation before the scroll decides your fate, and a busy, multi-element frame spends that fixation on confusion. If a reader can't say what the image is *of* in half a second on a phone, the image has already lost. A single expressive subject or a single moment of use, set in a real environment with enough texture to be interesting, sits naturally in the sweet spot — one focal point, real-world context around it. Composition depth is its own craft, and this series treats it right here in these opening documents; the default you carry everywhere is *one subject*.

There is one composition that breaks the single-subject default on purpose, and it's worth naming precisely because it's so often gotten wrong: the **side-by-side**. A side-by-side is not banned. It is a demanding-but-allowed composition — effectively two ad images sharing one frame, where each half gets half the pixels and the reader still gets only one first fixation. Because it asks so much of that single glance, it has to clear a higher bar, and every part of the bar is non-negotiable:

- **Per-half legibility.** Each half needs its own single, instantly legible focal subject. A busy half kills the entire frame, not just its own side.
- **One idea.** The two halves must express *one* idea or one clean contrast — not two separate ideas competing for the same glance. The reader should get the point of the pairing in that same half-second.
- **A directional cue.** Give the eye a reading order. A gaze, a lean, an implied motion that leads from one half to the other turns two competing subjects into one guided read — this is aimed gaze doing its attention-transfer job across the seam of the frame.
- **The phone-size test, as one unit.** Shrink the whole thing to feed scale and glance at it as a single object. If either half muddies, drop back to a single subject. No amount of

full-size polish rescues a side-by-side that falls apart small.

And one caution that lives beyond composition: a side-by-side that contrasts two states can start to read as a before/after result. The before/after ban is triggered by *content* — imagery presenting a bodily, health, or financial result as caused by the product — not by the two-panel layout itself. But the closer your state-contrast drifts toward that kind of transformation, the higher the compliance caution it demands. Think of it as a gradient: a side-by-side of two tidy vs. cluttered desks is low-risk; a side-by-side inching toward a physical or financial "look what happened to me" is playing with fire. The bans themselves are laid out in full in *Document 7, The Seven Bans*.

Dialing the Devices Across the Funnel

The devices don't change across the three places, but which one you reach for — and how hard you lean on it — shifts with the reader's state on your route.

- **The ad — the face as scroll-stopper and qualifier.** Here the attention device does its hardest work. The subject should look like the reader's situation — the person, or the pet, the reader recognizes as their own — so the face both stops the scroll and quietly self-selects the right audience. The gaze should serve the visual question the ad image is opening — often an aimed gaze toward the thing the reader should wonder about, or a reaction to something not fully shown (that curiosity partnership is the whole of *Document 3, Bounded Curiosity*). Hand-led, first-person shots are especially strong here for mechanism-led angles, where the doing *is* the story.
- **The advertorial page — the subject carries the narrative.** Past the click, faces stop being pure scroll-stoppers and start carrying the story. The same person, or the same kind of person, continues the narrative the ad image started — the reader who related to a face in the feed meets that face again on the page, and the thread holds.
- **Jump page (bridge page) — a face is optional.** This page is a second ad: the reader arrived on one click, still mostly cold, and the page's job is to get the video playing. A face here is a supporting player, not the lead. When you use one, direct gaze plus a credible, composed expression supports the direct promise the headline is making. Overacting here reads as desperation. Hand-led and first-person shots, meanwhile, work on all three — the doing translates everywhere.

Same toolkit, three settings: a face that stops and qualifies, a face that narrates, a face that reassures — with the moment of use ready to carry the story on any of them.

✓ Key Takeaways

- Expressive faces — human or animal — are pre-attentive: they win the first glance before conscious thought, which makes them the most reliable scroll-stopper and the most commonly wasted one.
- The gaze rule decides everything: **direct gaze** builds connection, **aimed gaze** transfers attention to the thing the reader should end on, and **jobless gaze** captures the eye and throws it away. Never use a face without a job.
- Expression is a claim; it must match the headline's angle, and natural emotion beats exaggerated, theatrical faces that readers are most likely to misread.
- At feed size: crop to the focal point — how tight scales with where the image lives — use one subject with one emotional read, and choose real over posed. Never use an identifiable person you can't use, and never a celebrity.
- The moment of use — the subject mid-action, including hands and first-person POV — is the second attention device, a full peer of the face: it demonstrates instead of displays, reads as UGC, and puts the reader inside the scene. Choose a face for feeling, the moment of use for doing.
- One focal subject is the default. The side-by-side is demanding-but-allowed: per-half legibility, one idea, a directional cue, and the phone-size test as one unit — with rising compliance caution as a state-contrast nears a bodily, health, or financial transformation (see *Document 7, The Seven Bans*).