

03

Bounded Curiosity: Show Less, Earn the Click

Document 3 of 8

CONTENTS

1 The Image Can Ask a Question

2 The Partial Reveal, Done Right

3 The Positive-Inference Rule

4 The Boundary: Bounded Curiosity vs. Bait

5 Curiosity That Qualifies

6 Calibrating Curiosity Down the Funnel

7 Key Takeaways

The Image Can Ask a Question

Curiosity usually gets taught as a headline device — the withheld how, the named mechanism, the open loop the reader clicks to close (that's the copy side, covered in *Document 3 of Copywriting Foundations, Curiosity: Withholding the How*). But the image can open a loop of its own, entirely without words, and controlled experiments confirm it pulls. A picture can pose a question, and the click is how the reader reaches for the answer.

A curiosity gap is the tension a person feels between what they can see and what they can't. In imagery, you create that gap by showing enough to establish that something interesting is present — and withholding the one piece that would resolve it. The object cropped so half of it sits out of frame. The reaction shown but not its cause. Hands mid-action on something you can't quite make out. The moment just *before* the interesting thing happens, or just *after*. In every case the reader's mind reaches for the missing piece, and reaching is the beginning of the click. The image, like a good headline, is showing the *what* and quietly withholding the resolving detail — the visual cousin of withholding the how.

This document is about doing that on purpose and doing it safely, because visual curiosity has a bright line running through it. On one side is bounded curiosity, which fills your funnel with the right readers. On the other is bait, which fills your account with problems. The whole craft is staying on the right side of that line.

The Partial Reveal, Done Right

The mechanism is concealment, but concealment is not the same as hiding. Hiding leaves the reader with nothing. A partial reveal shows the silhouette and withholds the workings — enough visible that the reader knows something worth knowing is there, not so much that the question answers itself.

Picture a fictional example: a countertop water device, and you want the ad image to open a loop. The bad version hides the whole thing in shadow — mysterious, sure, but about *what*? The good version shows a hand tilting a glass under the spout, the water caught mid-pour catching the light in a way ordinary tap water doesn't quite — and crops out the

device itself. Now the reader can see something is happening, can see it looks *good*, and is missing exactly one thing: what is that, and why does the water look like that? That's a partial reveal. The silhouette is legible; the workings are held back.

The concealment devices that survive at feed size are a small, dependable set. **Partial framing** — the key object cropped at the edge of the frame. **Occlusion** — hands or another object covering the crucial part. **Reaction-first** — an expressive reaction to something happening off-frame — a face, or any subject visibly reacting — so the reader wants to see what the subject is seeing. **Mid-process** — the moment before the result, the pour before the glass is full, the seal before it's pressed. Each of these shows the reader that a payoff is coming without handing it over.

Two structural rules keep a partial reveal from collapsing into noise. First, **the concealed element must be THE subject — singular**. Concealing a detail inside a busy, crowded image conceals nothing, because the reader never noticed the detail in the first place. This is exactly why the one-focal-subject default matters so much: curiosity needs a clear subject to build the gap around, and a cluttered frame has no gap because it has no focus. Second, **the image and the headline must open ONE question together, not two**. An image asking one thing while the headline asks another halves the pull of both — the reader's mind can't chase two gaps at once, so it chases neither. Expressive subjects are natural partners here: a visible reaction to something not fully shown is the classic feed-native curiosity image — the reaction-first device and a working expression in one shot.

The Positive-Inference Rule

Here is the finding that anchors this entire document, and it's the one most people miss. Partially concealing the key object raises engagement and preference **only when the reader's inference about the hidden part is positive**. Concealment is not a magic ingredient you sprinkle on any image to make it more interesting. It's an amplifier — and it amplifies whatever the reader *guesses* is behind it.

If what's implied reads as appealing — a solution, a satisfying result, something the reader already suspects they'd like — then concealment cranks up the pull. The reader's mind fills the gap with something good and wants to confirm it. But if the hidden part reads as confusing, alarming, or like nothing at all, concealment backfires completely. The reader feels lost instead of curious and scrolls on. This is the distinction to burn in: **ambiguity is**

not curiosity. A confusing image is just a confusing image. Mystery for its own sake — the shapeless "you won't believe this" of the visual world — leaves the reader with nothing to be curious *about*, which is the same failure the copy side calls out about vague headlines.

There's a clean test for it, and it takes five seconds. Describe what a cold reader — someone who has never seen your offer — would *guess* is behind the concealment. If the guess is specific and attractive ("that's got to be some kind of filtered water, and it looks better than my tap"), the concealment is working. If the honest answer is "I genuinely have no idea," you don't have a curiosity image. You have a confusing one, and you should start over. Run that test on every concealment you build.

The Boundary: Bounded Curiosity vs. Bait

Now the bright line. Every visual question you open is a check, and the next page in your funnel has to cash it.

Bounded curiosity poses exactly the question the funnel goes on to resolve. The reader clicks the ad, lands on the advertorial, sees the same subject continued, and gets the answer as the story unfolds. The scent stays intact from image to page — the water device the ad image teased is the water device the advertorial explains. Nothing is betrayed. The reader arrives already leaning in, because the page is delivering precisely what the image promised. That's a click that converts.

Bait is a visual question the funnel never answers. A shocking implied reveal that never materializes. A subject with nothing to do with the offer, stapled on because it grabs eyeballs. Drama borrowed from nowhere and paid back never. Bait absolutely produces clicks — that's the seductive part, and it's why people keep reaching for it. But it produces clicks and nothing else: the reader lands, finds the promise wasn't real, and bounces.

The reason this matters so much more than one bad campaign is what the bounce does to you downstream. The reader's confusion isn't just a lost sale — it's a signal. The native network's quality systems watch what happens after the click, and a wave of readers who bounce instantly tells those systems your creative is misleading. The account pays for that in throttled delivery and rising costs, and not only on the one ad that baited them. You are teaching the network's machinery that your traffic is low-quality, and that judgment follows you. Beyond the economics, misleading and irrelevant images are flat-out banned by every

native network — bait is a policy violation before it is ever a strategy mistake (see *Document 7, The Seven Bans*). It erodes network trust, not just a single campaign's numbers.

Curiosity That Qualifies

The best visual curiosity does two jobs at once: it pulls the right reader in *and* lets the wrong reader scroll past. A curiosity image built out of the target reader's own world — their kind of kitchen, their kind of ache, their kind of clutter — reads as "wait, what is this about?" to the person it's for, and as pure background noise to everyone else. That self-selection is more valuable than raw click volume, because a click from the wrong reader is just the front end of a bounce. Concealment aimed at a specific reader's situation fills the funnel with people the advertorial can actually move — which is the same thing bounded curiosity is really for. You're not trying to make everyone curious. You're trying to make the *right* everyone curious, and let the rest keep scrolling.

That reframes the whole point of an image loop. Bounded curiosity is a promise, and the next page has to pay it off. Show less in the ad image, yes — but only ever less of something the reader will be glad to find on the other side of the click. The gap you open is a debt the funnel takes on. Keep it small enough, specific enough, and honest enough that the very next page can settle it, and the reader converts. Open a gap the funnel can't close, and you haven't earned a click — you've written a check that bounces, twice: once for the reader, and once with the network watching.

■ Calibrating Curiosity Down the Funnel

Visual curiosity, like the copy kind, is front-loaded at the ad — but how it gets spent depends on which route your funnel runs. The advertorial route pays the gap off through a story; the jump-page route keeps it alive and hands the payoff to the video.

- **The ad — the gap at full width.** This is curiosity's home. The image poses the visual question — partial framing, occlusion, a reaction to something off-frame — and resolves nothing. The whole point of the ad image is to open the loop the click will reach to close, and to open it out of the target reader's world so it qualifies as it pulls.

- **The advertorial page — the gap starts paying off.** Past the click, the same subject the ad image teased continues, and the answer arrives as the story unfolds. The concealment from the ad image is honored here: the reader finally sees the device, understands the moment, gets the "what" the image promised. Scent intact, question answered. The page can open smaller next-questions to keep the reader moving, but it must first settle the one the image opened.
- **Jump page (bridge page) — the gap stays open and points at the video.** This reader clicked the ad and landed here directly — the loop the ad image opened is still almost fully open, and the video is where it gets closed. So the jump-page image doesn't resolve the mystery on the page; it restates it and aims it at the play button: *watch how it actually works*. What it must not do is open a *different* gap — a fresh, unrelated visual mystery breaks the scent the click was following. Same loop, carried forward, handed to the video.

Same gap, two routes: on the advertorial route it closes through the story; on the jump-page route it stays open and gets handed to the video.

✓ Key Takeaways

- The image itself can open a curiosity gap — showing enough to prove something interesting is present while withholding the piece that resolves it. It's the visual cousin of withholding the how.
- A **partial reveal** shows the silhouette and hides the workings, using devices that survive feed size: partial framing, occlusion, reaction-first, and mid-process.
- The **positive-inference rule** governs everything: concealment only helps when the reader's guess about the hidden part is *good*. Ambiguity is not curiosity — if a cold reader's guess is "I have no idea," start over.
- The concealed element must be the single subject, and the image and headline must open one question together, not two.
- **Bounded curiosity** poses exactly the question the funnel answers, keeping scent intact and converting the click. **Bait** teases what the funnel never delivers — and it

damages account delivery and network trust far beyond one campaign, on top of being an outright policy violation (see *Document 7, The Seven Bans*).

- The best curiosity qualifies: built from the target reader's world, it pulls the right reader in and lets the wrong one scroll past. Every gap you open is a promise the next page must pay off.