

07

The Seven Bans: Images That End Campaigns

Document 7 of 8

CONTENTS

1 The Image Speaks for Itself

2 Ban One: Before/After Transformation

3 Ban Two: Sexualized Imagery and Body-Part Emphasis

4 Ban Three: Celebrities and Public Figures

5 Ban Four: Fake UI Elements

6 Ban Five: Shock, Gore, and Disturbing Imagery

7 Ban Six: Misleading or Irrelevant Images

8 Ban Seven: Implied Disease and Medical Claims

9 Where the Bans Bite: Ad Images First

10 Key Takeaways

The Image Speaks for Itself

Everything else in this series is about making an image *work*. This document is about making sure it doesn't get you thrown out of the room before it gets the chance.

Here is the frame that governs all of it. Regulators and networks don't judge your ad by the text alone, or the image alone, or the fine print alone. They judge it by its **net impression** — the single message a reasonable person walks away with after taking in every element at once. And the part most media buyers underestimate is this: an image is fully capable of making a claim entirely on its own.

A photo that implies dramatic weight loss *is* a weight-loss claim. A photo that implies a disease was cured *is* a medical claim. It doesn't matter that your headline was careful, that your copy was hedged, that you never typed the forbidden sentence. "The text never said it" is not a defense — if the picture said it, your ad said it. The image is not decoration hanging off a compliant headline; it is a claim-making device in its own right, and it will be read as one.

Once you internalize that, the seven bans stop feeling like an arbitrary blacklist and start feeling like what they are: seven predictable ways an image can make a claim, or cross a line, that ends the campaign carrying it. These are not per-network technicalities to optimize around. Every major native network prohibits all of them, and they're taught here as universal. One non-compliant image can get an ad rejected, a campaign paused, or an account flagged — and account standing is the scarcest resource you own. Compliant images are unlimited; the account that runs them is not.

So we walk each ban the same way: what it is, why it bites, and what to do instead. The goal isn't to memorize a list of nos. It's to develop the reflex of looking at any candidate image and asking the reviewer's question automatically — *what claim does this picture make at a glance?*

Ban One: Before/After Transformation

The single most-wanted image in health and beauty is the one you cannot run.

Transformation photos — the visible weight loss, the cleared skin, the regrown hair, the whitened teeth, the fattened bank account — that show the change *itself* are explicitly banned across the networks. And the ban follows you into implied territory too:

"transformation" collages, dramatic same-person body contrasts, the split-frame that puts a struggling state on one side and a triumphant one on the other.

Here is the part that trips people up, so read it slowly: **the ban is about content, not layout.** What triggers it is imagery that presents a bodily or financial result as caused by your product. A side-by-side arrangement is not itself illegal — putting two images beside each other is a composition choice, and a demanding one, covered in its own right (see *Document 2, The Attention Devices*). What gets you rejected is not the two panels; it's that the two panels together say *"this product transformed this body."* You could commit the same violation in a single frame, and you could arrange a compliant side-by-side that never touches a transformation claim. The reviewer isn't counting your panels. They're reading your net impression.

What to do instead: **show one state, and let the words carry the change over time.**

Show the desirable end state — the person already living the good outcome — or show the relatable problem state that your reader recognizes as their own. Pick the state that serves the angle, render it as a single honest moment, and let the headline do the work of implying a *before* and an *after*. The change lives in the reader's head, sparked by the copy. It never lives, provably and product-caused, in the picture.

Ban Two: Sexualized Imagery and Body-Part Emphasis

No sexualized framing. No zoomed or cropped emphasis on intimate or isolated body parts — the midriff, the chest, the close-cropped patch of skin. No imagery whose composition is engineered to make a reader self-conscious about their own body.

Health and fitness angles are where good marketers walk straight into this one, because the "belly fat" zoom feels like it's just illustrating the problem. It isn't. That crop is a ban on two counts at once: it's body-part emphasis, and it's an implied claim about what the

product does to that body part. The reviewer sees a picture engineered to make someone feel bad about their midsection, attached to a product that promises to fix it — that's the whole banned pattern in one frame.

What to do instead: **whole people, in context, doing things**. A person is not a body part. Someone photographed as a complete human being — in a real setting, mid-action, with an expression — carries all the relevance you need without isolating and objectifying a slice of anatomy. If your angle is about the body, show the person the body belongs to, living their life.

Ban Three: Celebrities and Public Figures

No celebrities. No politicians, no famous doctors, no recognizable scientists, no public figure whose face a scroller would know — not in ad imagery, period. An implied endorsement by a recognizable face is one of the fastest routes to an account ban that exists, because it's the kind of violation that draws complaints from the person depicted, not just the platform.

Two extensions worth stating plainly. This covers people who are *near-famous* enough to be mistaken for the real thing, so the borrowed-credibility instinct — "a picture that reads like a well-known TV doctor" — is exactly the instinct to kill. And it covers **AI-generated likenesses of real people**, which are equally banned and, increasingly, outright illegal. Generating a face that resembles a specific celebrity is not a clever workaround; it's the same violation with an extra layer of risk stacked on top.

What to do instead: **unknown faces**. Ordinary-looking people and AI-generated humans who resemble no one in particular — the whole point of the attention-device work in this series (see *Document 2, The Attention Devices*) is that a relatable stranger does the attention job *better* than a borrowed celebrity, and does it without betting your account on it.

Ban Four: Fake UI Elements

No fake play buttons. No fake close buttons or X's, no fake system dialogs, no fake notification badges, no graphic pretending to be a functional control. The governing rule is simple: an image must not appear to have functionality it doesn't have.

Every network treats this as deliberate deception, and they're right to — a fake play button is a lie told to a finger. The reader taps expecting video and gets a static image and a redirect. That's not a gray area or a matter of taste; it's the kind of thing that gets an account actioned fast, because it's unambiguous intent to trick.

What to do instead: **imply motion photographically, never graphically.** If the appeal of your creative is that something is happening — a moment of use, a device mid-action, a scene caught in progress — sell that with the photograph itself: a paused-moment feel, a gesture frozen mid-motion, a composition that reads as a still lifted from a real video. You can absolutely make an image feel alive. You just can't bolt a fake control onto it.

Ban Five: Shock, Gore, and Disturbing Imagery

No blood, no wounds, no rashes or skin conditions, no medical procedures, no dead bodies, no bodily fluids, no gross-out content — and this holds *even when the offer is health-related and the condition is genuinely real*. "But that's what the problem actually looks like" is not an exception. It's the exact case the ban exists for.

The reasoning is both compliance and craft. Shock stops the scroll, yes — and then it kills the account, because networks read gore as a serving-quality violation regardless of context. But it also fails on its own terms: a disturbing image repels far more readers than it recruits, and the ones it stops aren't the ones who convert.

What to do instead: **represent the problem through the person's experience, not through graphic depiction.** The discomfort of a skin condition can live in a person's expression and posture — the self-conscious tug at a sleeve, the tired face — without ever showing the condition itself. The reader recognizes the *feeling*, which is more persuasive than the symptom anyway. Symptom-level, experience-level framing is the ceiling here, and note that even that ceiling stays underneath this ban's limits: you evoke the experience, you never depict the gore.

Ban Six: Misleading or Irrelevant Images

The image has to be an accurate representation of what the funnel is actually about. No borrowed drama. No unrelated eye-candy stapled on because it gets clicks. No misleading crop or zoom that makes something appear to be what it isn't. If the only connection

between your image and your offer is that the image is arresting, you're in violation.

This ban is enforced twice over, which is what makes it worth respecting even when you think you could sneak it past review. The networks enforce it directly — an image that misrepresents the offer is a policy breach. And the funnel enforces it again, mechanically and without appeal: a click bought with a false promise doesn't convert. The reader arrives, feels the mismatch between the picture and the page, and leaves. You paid for attention that curiosity's own boundary rules already told you would evaporate (see *Document 3, Bounded Curiosity* on where the honest tease ends and bait begins).

What to do instead: **cast the image for the offer's real story**. Relevance isn't a compliance tax; it's what makes the click worth buying. The most accurate image and the most profitable image are usually the same image, for the same reason.

Ban Seven: Implied Disease and Medical Claims

No imagery that implies your product treats, cures, or reverses a disease or medical condition. No syringes-and-cure visuals. No disease-name props sitting in frame. No clinical staging — the white coat, the exam table, the medical device — that implies a medical intervention is taking place. No imagery suggesting a supplement steps in for actual medication.

This is Ban One's cousin: both are cases where the picture makes a claim the words were careful never to make. A bottle of supplements photographed beside a syringe and a stethoscope has said "this is medicine that treats disease" without a single word of copy. The net impression is a medical claim, and a medical claim from a non-medical product is one of the brightest lines in advertising.

What to do instead: **stay at the symptom and experience level** — how a person feels, how their day goes, the discomfort they live with — and keep even that inside Ban Five's limits on depiction. You can show a person who sleeps better; you cannot show the disease you're implying you cured. The ceiling is the lived experience. The floor drops out the moment the image starts playing doctor.

Where the Bans Bite: Ad Images First

Here is the calibration that matters for this document — not how the concept dials down the funnel, but where the *enforcement* lands hardest, because that changes how you spend your caution.

The bans are enforced most aggressively on your **ad images**: the image the network actually reviews and serves into the feed. That creative goes through a review gate, gets served at scale, and draws the most scrutiny, so it is where a violation is most likely to be caught and most expensive when it is. Screen your ad images against all seven bans with zero tolerance. Near a line is over a line. When you're weighing whether an ad image might be a problem, the answer is to pick a different image — the supply is unlimited and the account is not.

Your landing pages — the advertorial and the jump page (bridge page) — face *less* scrutiny on some elements. A picture that would get an ad instantly rejected sometimes lives on a page further down the funnel without the same immediate consequence. It is tempting to read that as permission. It is not. That gap is a difference in **enforcement pressure, not in the rule**. The net-impression principle applies to the whole funnel — every page a reader touches contributes to the single message the ad is judged to have made — so a banned claim buried on your jump page is still a banned claim, and it can still emerge in a review, a complaint, or an audit and take the whole funnel down with it.

So the working posture is: treat the seven bans as absolute on your ad image, and carry the same net-impression discipline onto your pages even where the scrutiny is lighter. The lighter scrutiny buys you a little room to breathe, not a license to break the rule. Caution still — just calibrated to where the enforcement actually lands.

✓ Key Takeaways

- Net impression governs everything: an image can make a claim entirely by itself, and "the text never said it" is no defense. If the picture said it, the ad said it.
- The before/after ban is triggered by content, not layout — imagery presenting a bodily or financial result as product-caused. Side-by-side is a composition choice

(see *Document 2*), not the violation. Show one state and let the headline carry the change.

- Body-part emphasis and sexualized framing: show whole people in context doing things, never isolated anatomy.
- No recognizable public figures — including AI likenesses of real people. Use unknown, ordinary faces or AI-generated people who resemble no one.
- No fake functional UI, no shock or gore (even when the condition is real), no misleading or irrelevant images, and no imagery implying disease treatment or cure. Stay at the symptom-and-experience level.
- Screen every candidate against all seven bans before you fall in love with it, and remember the enforcement reality: the bans bite hardest on ad images, while landing pages face lighter scrutiny on some elements — but the same net-impression principle applies to the whole funnel, so caution still.