

08

From Angle to Image: The Working Process

Document 8 of 8

CONTENTS

1 Seven Documents, One Motion

2 Step 1 — Start From the Angle, Never From Browsing Images

3 Step 2 — Decide Where the Image Lives

4 Step 3 — Choose the Subject

5 Step 4 — Decide What to Show and What to Conceal

6 Step 5 — Generate or Source in the UGC Register

7 Step 6 — Run the Glance Test

8 Step 7 — Screen Against the Seven Bans, Then Stop

9 Key Takeaways

Seven Documents, One Motion

You now have the whole toolkit. You know what each of the three places asks an image to do, how an ad image earns a glance, how expressive subjects and the moment of use pull attention, where honest curiosity ends and bait begins, why UGC beats polish, how color works against the feed, and which seven images will end a campaign. That's a lot of doctrine to carry into a moment that, in practice, lasts about ninety seconds: *"I need an image for this ad — now what?"*

This document is the answer to that question. It introduces nothing new. Every rule here already lives in one of the seven documents before it; what's been missing is the *order*. A rule you know but reach for at the wrong moment is nearly as useless as a rule you don't know at all — pick the image before you've settled the angle and no amount of craft downstream will save the creative.

So think of this as the assembly line the previous seven docs feed into. Seven steps, in sequence, each one narrowing the field the next step works on. Run them in order and "how do I pick an image?" stops being a moment of paralysis and becomes a process you can trust — and hand to anyone. The steps *are* the structure here; there's no separate calibration section, because the workflow itself stops to ask where the image lives and calibrates as it goes.

Step 1 — Start From the Angle, Never From Browsing Images

Image selection begins before any image exists. It begins with the angle.

This is the discipline that separates working creative from pretty wallpaper, and it's the one most often skipped, because browsing images *feels* like progress. It isn't. Choosing an image first and retrofitting a story around it produces creative that looks good and communicates nothing — an image doing a job it was never cast for. The angle is what decides what the image must say; the image is then cast for that role the way you'd cast an actor for a part that's already written.

So the rule is blunt: if the angle isn't settled, stop. You are not ready to look at images yet. Settle the angle — decide the single idea this creative is arguing — and only then let it tell you what the image needs to communicate. Every step that follows is downstream of this one, which is why it's first. (The angle-to-image casting logic lives in *Document 6, Congruence: One Idea Through the Whole Funnel*.)

Step 2 — Decide Where the Image Lives

The same image can be exactly right in one place and flatly wrong in another. Before you judge a single candidate, decide where this image is going to live.

There are three places it can live — spread across the two funnel routes, ad → advertorial → sales page or ad → jump page → VSL — and each asks a different job of it. The **ad image** has to win qualified attention from a cold scroller who owes you nothing. The **advertorial image** has to confirm the scent the click was chasing — to reassure a reader who just arrived that they're in the right place. The **jump page (bridge page) image** has to support the promise with credibility for a reader who is only one click past the feed — a second ad's image, earning the play. Three different readers, three different jobs.

This is where the workflow names the destination, so it's worth naming what's riding on it: every later judgment in this process — which subject, how much polish, how assertive the image is allowed to be — calibrates to that place's job and the reader's actual state there. Get this wrong here and you'll make correct-looking decisions in service of the wrong goal. (The full breakdown of each job lives in *Document 1, The Three Jobs of an Ad Image*.)

Step 3 — Choose the Subject

Now decide how a subject actually shows up in the frame. In practice this is a choice between two attention devices, and the choice is dictated by what your angle is about.

Reach for **an expressive face — human or animal — when the story is about the feeling**: their relief, their skepticism, their quiet satisfaction — or the dog's ears-up delight. And when you choose a face, immediately give its gaze a job, because a face with no job is wasted attention. A *direct gaze* looks out at the reader and does connection and

relatability — the "someone like me" handshake. An *aimed gaze* looks at something inside the frame and transfers the reader's attention to whatever the subject is looking at, so the reader ends where you want them to end.

Reach for **the moment of use when the story is about the doing** — the mechanism, the ritual, the moment itself: hands mid-task, first-person point of view (the reader's-own-hands framing that puts them inside the moment), or the subject itself mid-action. This device is a peer of the face, not a fallback. An object sitting alone is a product shot; the same object in a hand is a *moment*, and moments are what stop scrolls. If your angle is about how something is used or what using it feels like, the moment of use will almost always beat a face.

(The craft rules for both devices — crop to the focal point, one subject, one action, expression matched to the promise — live in *Document 2, The Attention Devices*.)

Step 4 — Decide What to Show and What to Conceal

With the subject and device chosen, decide how much the image reveals. This is the curiosity dial, and it's a step, not an afterthought.

The move you're looking for is a **bounded visual question**: partially conceal the object of curiosity while the visible part carries a clearly positive inference. Enough is shown that the reader forms a good expectation; enough is hidden that they have to click to resolve it. That gap, honestly built, earns the click that the next page then pays off.

But concealment only works when what's hidden genuinely promises something good — the tease has to be a promise the destination can keep. Hide everything and you've got a blank; tease something the page never delivers and you've built bait. Bait is not an aggressive tactic; it's an account-killer, because the mismatched click doesn't convert and the pattern gets you flagged. Conceal to *invite*, never to *trick*. (The full concealment boundary — where the honest question ends and bait begins — lives in *Document 3, Bounded Curiosity*.)

Step 5 — Generate or Source in the UGC Register

Now, and only now — angle set, destination chosen, subject and reveal decided — do you actually produce candidates. And you produce them in the default register: **UGC-style**.

That means candid, phone-shot feel, real settings, imperfect lighting — an image the surrounding editorial feed could plausibly contain, rather than one that announces itself as an ad. Obvious stock photography is the cardinal sin; it reads as commerce instantly and the cold scroller's guard goes up.

When you're generating with AI, prompt for that register in plain, ordinary language — ask for candid, documentary, phone-camera looks, the way you'd describe a real snapshot to a friend. Use negative prompts to push away from studio polish and glossy perfection. Generate many and shortlist few; volume is cheap and your standard should be high. And check every candidate for artifacts — anatomy above all, hands and faces especially — before it goes anywhere near a campaign, because a six-fingered hand or a melted background undoes all the authenticity you were reaching for. (The register doctrine, the AI prompting guidance, and the narrow late-funnel exception for slightly cleaner imagery all live in *Document 5, The UGC Default*.)

Step 6 — Run the Glance Test

Judge the shortlist at the size and setting where it will actually be seen: feed size, on a phone, in a feed. Not blown up on your monitor where every candidate looks competent. At real render size, most of them won't.

The bar is fast: one focal subject, instantly legible, no text in the image, and edges that hold against the page background rather than dissolving into it. A composition that needs two seconds to parse has already lost — the scroller was gone at one. If you can't tell what a candidate is about in the first glance, neither can your reader, and no amount of merit at full size will rescue it.

And if a candidate is a **side-by-side**, hold it to the demanding standard, because that composition is allowed but unforgiving: each half has to be legible on its own, the pairing has to express exactly one idea, a directional cue has to order the read so the eye knows which half comes first, and the whole thing has to pass this glance test as a single unit —

not two images sharing a frame, but one image made of two parts. Most side-by-sides fail one of those tests. Judge them here, honestly, before compliance ever sees them. (The side-by-side standard lives in *Document 2, The Attention Devices*.)

Step 7 — Screen Against the Seven Bans, Then Stop

The final gate, applied before you let yourself fall in love with any surviving candidate: the seven absolute bans.

Screen each one the way a network reviewer or a regulator would — by what the image communicates as a whole, its **net impression**, not by its layout or your intent. Ask the reviewer's question: *what claim does this picture make at a glance?* The bans are content-triggered, so a candidate that feels innocent to you can still make a banned claim to a stranger seeing it cold. Enforcement is hardest on your ad images; your landing pages get less scrutiny on some elements, but the same net-impression principle applies across the whole funnel, so caution still. (The seven bans, each in full, live in *Document 7, The Seven Bans*.)

An image that fails here fails completely — no matter how brilliantly it did on every earlier step. A perfect subject, a flawless glance test, an inspired concealment: all of it is void the instant the image makes a banned claim. Compliance isn't the last item on a list of equals; it's the gate the whole list has to pass through.

And an image that passes is **done**. That's the last step. Chosen, checked, and ready for the place it lives — the workflow ends here. There is no step eight. You don't loop back, you don't add a stage, you don't keep tinkering. Seven steps in, you have an image you can trust, and trusting the process means knowing when it's finished. Ship it and move on to the next angle — which, because you now know how, starts the seven steps over again.

✓ Key Takeaways

- The process runs in one direction: angle → destination → subject → reveal → generate → glance test → compliance. Each step narrows the field the next one works on, so running them in order is the whole point.

- Always start from the angle, never from browsing images — the angle decides what the image must say, and the image is cast for that role. If the angle isn't settled, stop.
- Choose the subject by what the story is about: an expressive face — human or animal — with a gaze given a job (direct for connection, aimed to transfer attention) when it's about the feeling, or the moment of use — hands, first-person POV, or the subject mid-action — when it's about the doing.
- Produce candidates in the UGC register — candid, phone-shot, feed-plausible — prompting AI in plain language, generating many, shortlisting few, and checking anatomy for artifacts.
- Judge the shortlist at feed size on a phone; a composition that takes two seconds to parse has lost, and any side-by-side must clear the demanding standard as a single legible unit.
- Screen every survivor against the seven bans by net impression, then stop. An image that passes is chosen, checked, and done — there is no step after compliance.